

# Deep Learning

# Deep learning

Representation learning with a hierarchy of concepts

Those concepts are represented by layers in a neural network model

## Unsupervised models

- Autoencoders

## Supervised models

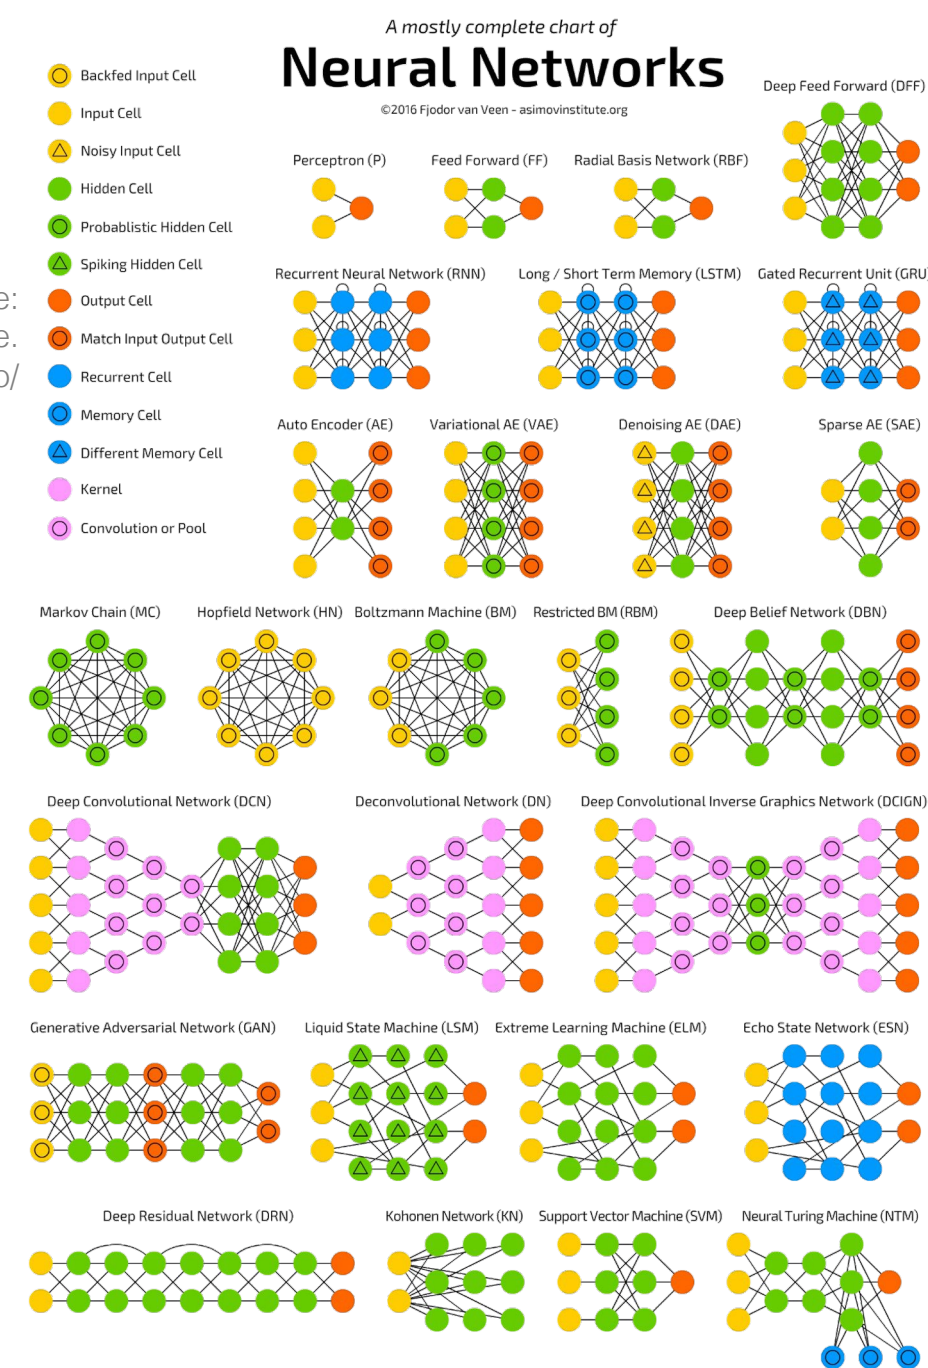
- Image analysis:
  - Convolutional Neural Networks
- Timeseries analysis (often NLP):
  - Recurrent Neural Networks (e.g. LSTMs)
  - Transformer models

## Generative models

- Generative Adversarial Networks (GANs)
- Diffusion Models (e.g. DALL-E 2, Stable Diffusion)
- Generative Pre-trained Transformer (GPT)

# Types of Deep Learning Tools

Azimov Institute:  
<http://www.asimovinstitute.org/neural-network-zoo/>



# Autoencoders


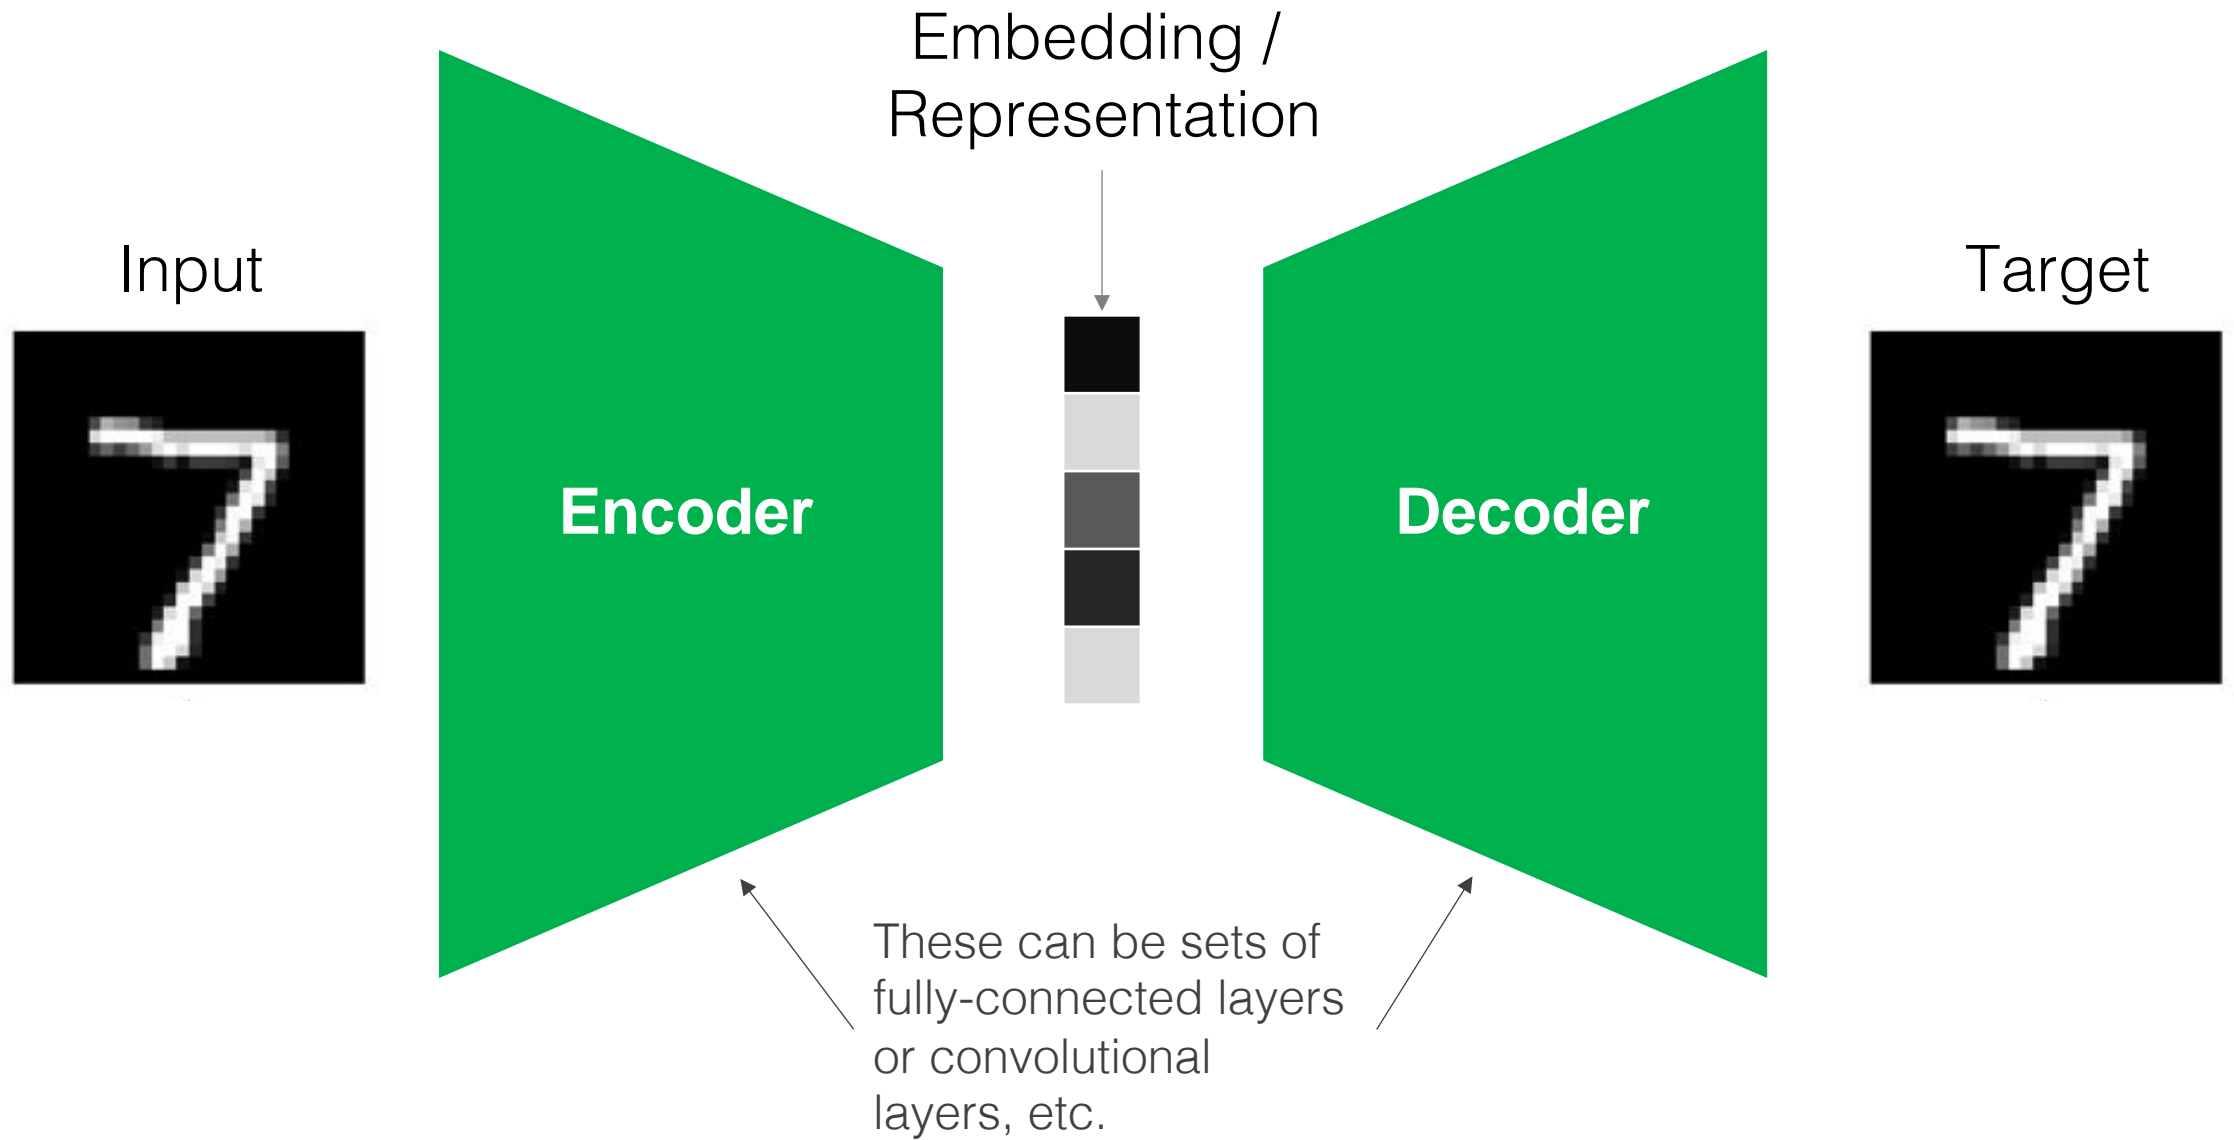


Image from: <https://www.jeremyjordan.me/autoencoders/>


# Autoencoders

Our goal is often to develop a good **encoder** that represents our features well




# Convolutional Neural Networks


# AlexNet



Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.



## Key





# Convolutional Neural Networks

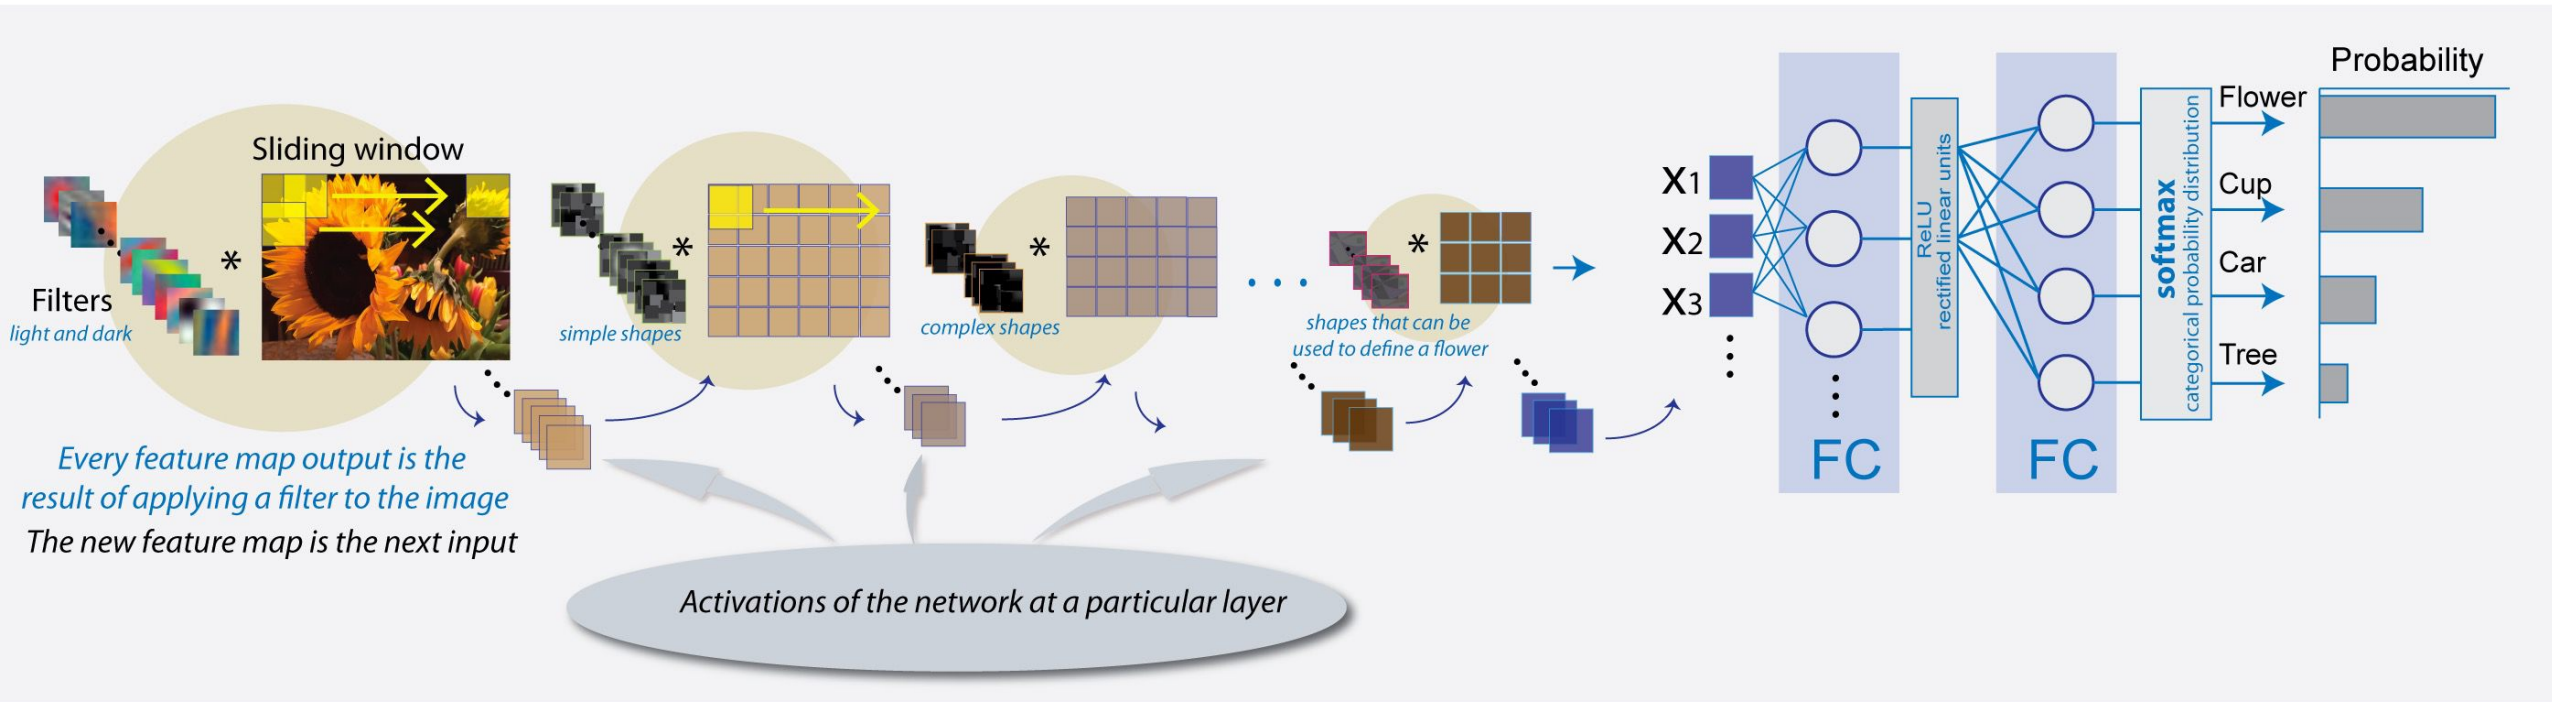
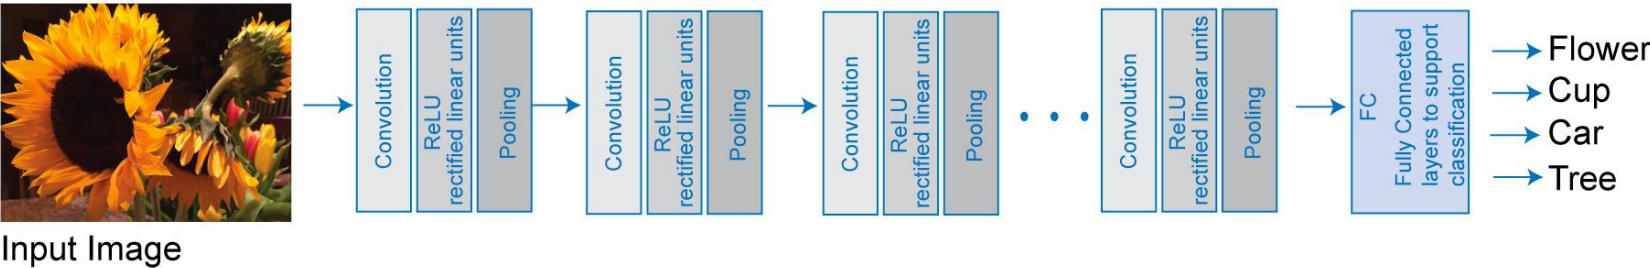


Image from the Mathworks



Data:  $\mathbf{x}$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6



Weights:  $\mathbf{w}$

1	1	1
0	0	0
-1	-1	-1

=

Output:  $\mathbf{x} * \mathbf{w}$


# 2D Convolution

Data:  $x$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $x * w$


Computing  
one output  
value:

$$1 \cdot 1 + 1 \cdot 2 + 1 \cdot 5 +$$

$$0 \cdot 0 + 0 \cdot 2 + 0 \cdot 3 +$$

$$(-1) \cdot 4 + (-1) \cdot 5 + (-1) \cdot 5$$

# 2D Convolution

Data:  $x$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $x * w$

-6			

Computing  
one output  
value:

$$\begin{aligned}
 &1 \cdot 1 + 1 \cdot 2 + 1 \cdot 5 + \\
 &0 \cdot 0 + 0 \cdot 2 + 0 \cdot 3 + \\
 &(-1) \cdot 4 + (-1) \cdot 5 + (-1) \cdot 5 = -6
 \end{aligned}$$

# 2D Convolution

Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $X * w$

-6	-11		

Computing  
one output  
value:

$$\begin{aligned}
 &1 \cdot 2 + 1 \cdot 5 + 1 \cdot 1 + \\
 &0 \cdot 2 + 0 \cdot 3 + 0 \cdot 2 +
 \end{aligned}$$

$$(-1) \cdot 5 + (-1) \cdot 5 + (-1) \cdot 9 = -11$$

# 2D Convolution

Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $X * w$

-6	-11	-12	

Computing  
one output  
value:

$$\begin{aligned}
 &1 \cdot 5 + 1 \cdot 1 + 1 \cdot 4 + \\
 &0 \cdot 3 + 0 \cdot 2 + 0 \cdot 0 + \\
 &(-1) \cdot 5 + (-1) \cdot 9 + (-1) \cdot 8 = -12
 \end{aligned}$$

# 2D Convolution

Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $X * w$

-6	-11	-12	-11

Computing  
one output  
value:

$$\begin{aligned}
 &1 \cdot 1 + 1 \cdot 4 + 1 \cdot 2 + \\
 &0 \cdot 2 + 0 \cdot 0 + 0 \cdot 0 +
 \end{aligned}$$

$$(-1) \cdot 9 + (-1) \cdot 8 + (-1) \cdot 1 = -11$$

# 2D Convolution

Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $X * w$

-6	-11	-12	-11
-7			

Computing one output value:

$$1 \cdot 0 + 1 \cdot 2 + 1 \cdot 3 +$$

$$0 \cdot 4 + 0 \cdot 5 + 0 \cdot 5 +$$

$$(-1) \cdot 6 + (-1) \cdot 3 + (-1) \cdot 4 = -7$$

# 2D Convolution



Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

6 x 6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

3 x 3

\*

=

Output:  $X * w$

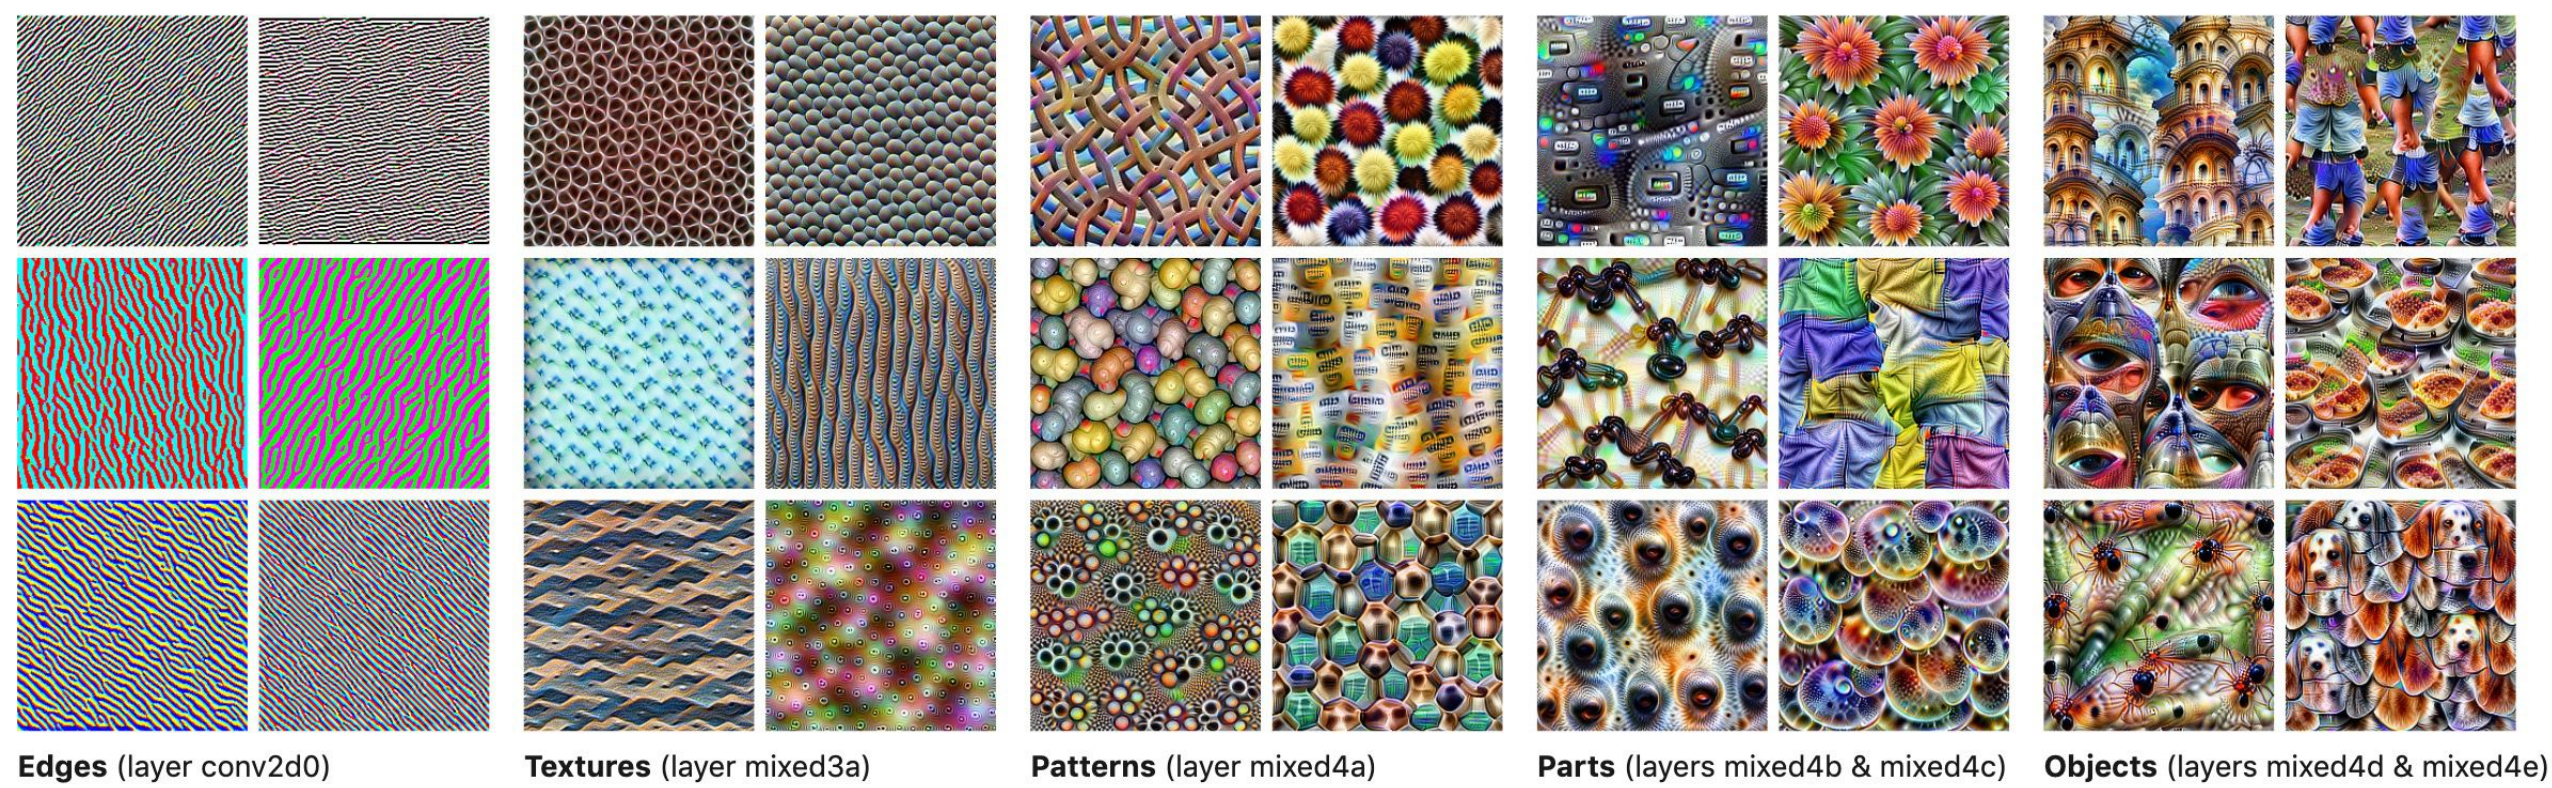

<b>-6</b>	<b>-11</b>	<b>-12</b>	<b>-11</b>
<b>-7</b>	<b>-2</b>	<b>-2</b>	<b>-4</b>
<b>4</b>	<b>1</b>	<b>-2</b>	<b>1</b>
<b>3</b>	<b>-4</b>	<b>-6</b>	<b>-10</b>

4 x 4

# 2D Convolution



# What features do layers respond to?



Olah et al, 2017: <https://distill.pub/2017/feature-visualization/>



# Features


**Dataset Examples** show us what neurons respond to in practice



**Optimization** isolates the causes of behavior from mere correlations. A neuron may not be detecting what you initially thought.



Baseball—or stripes?  
*mixed4a, Unit 6*



Animal faces—or snouts?  
*mixed4a, Unit 240*



Clouds—or fluffiness?  
*mixed4a, Unit 453*



Buildings—or sky?  
*mixed4a, Unit 492*

Olah et al, 2017: <https://distill.pub/2017/feature-visualization/>


# Resources on Visualization of Features

Feature visualization: <https://distill.pub/2017/feature-visualization/>

Building blocks of interpretability: <https://distill.pub/2018/building-blocks/>


Activation Activation Atlases: <https://distill.pub/2019/activation-atlas/>

# Convolution Layer



From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018


For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!

From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions




$$\text{Parameters} = (5 \times 5 \times 3) \times 6 = 450$$

$$(5 \times 5 \times 6) \times 10 = 1,500$$

From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018




# 1 x 1 Convolution Explained




1x1 CONV  
with 32 filters

(each filter has size  
1x1x64, and performs a  
64-dimensional dot  
product)




From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

# Max Pooling




max pool with 2x2 filters  
and stride 2




From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

# Fully Connected Layer



# AlexNet




Input
11x11 conv, 96
5x5 conv, 256
max pool
3x3 conv, 384
max pool
3x3 conv, 384
3x3 conv, 256
max pool
FC 4096
FC 4096
FC 1000
softmax

Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.


## Key

Input or output layer
Convolutional Layer
Fully Connected Layer
max pooling layer


AlexNet  
(2012)



VGG16  
(2014)



VGG19  
(2014)



Note: an activation function is applied to the output of each layer

Fewer layers,  
larger filters

## Key

Input or output layer

Convolutional Layer


Fully Connected Layer

max pooling layer







# CNN Architectures

Adapted from Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

# Residual Networks (ResNet)



Skip Connection enable faster convergence, more effectively backpropagate the error signal (avoiding vanishing gradients)




He, K., Zhang, X., Ren, S. and Sun, J., 2016. Deep residual learning for image recognition. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 770-778).





Scene  
classification



Object  
detection

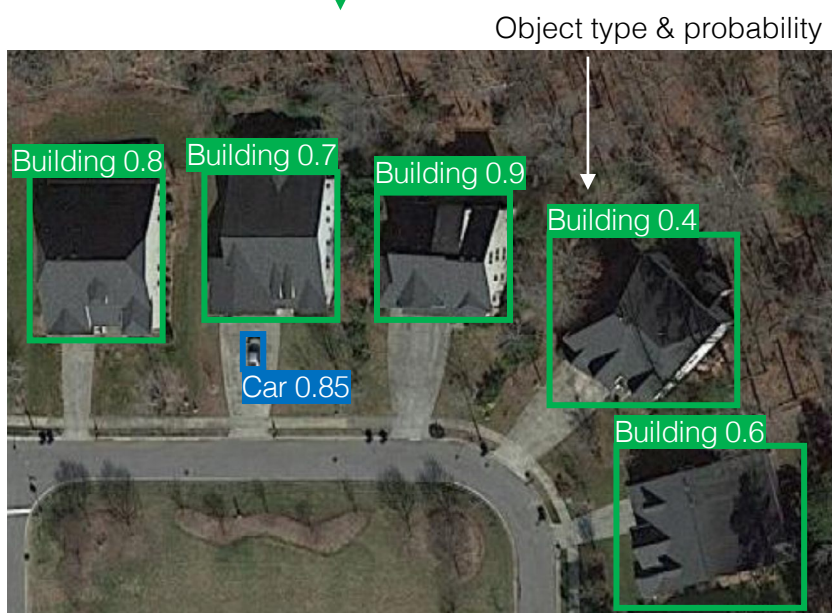


Image  
segmentation






## Scene classification



AlexNet  
VGG  
GoogLeNet  
ResNet

Inception  
DenseNet  
SqueezeNet  
EfficientNet

Confidence or probability

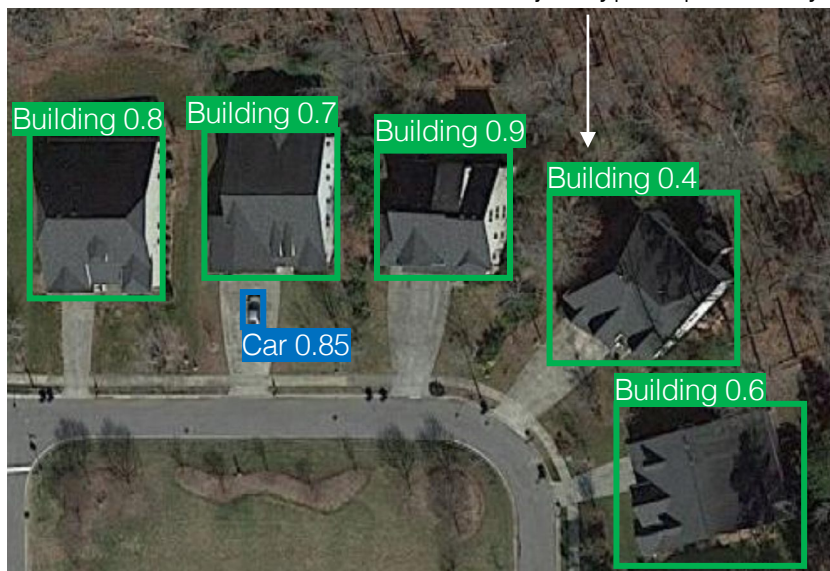


## Object detection



Faster/Fast/R-CNN  
Mask R-CNN  
YOLO  
Single Shot Detector (SSD)  
RetinaNet

Object type & probability



## Image segmentation



U-Net (2015)  
SegNet (2016)  
DeepLab (2017)  
FCN (2016)





# ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

Fei-Fei Li et al. 2010 ([link](#))


Competition at:  
Conference on Computer Vision and  
Pattern Recognition (CVPR)

## USED FOR MODEL PRETRAINING

AlexNet  
(Hinton, Sutskever, and Krizhevsky)

Human error ~5%



perfect



Source: Quartz, [link](#) David Yanofsky | Quartz

Data: ImageNet

# Deep Learning Models Compared



Models compared for ImageNet


Many of these models are available through Keras ([link](#))

A. Canziani, E. Culurciello and A. Paszke, "Evaluation of neural network architectures for embedded systems," *2017 IEEE International Symposium on Circuits and Systems (ISCAS)*, Baltimore, MD, 2017, pp. 1-4.

# Deep learning frameworks


Tensorflow ([link](#))

Framework for implementing graphical models, such as neural networks



Keras ([link](#))

Wrapper for Tensorflow to make coding easier: higher level and excellent API



Keras

PyTorch ([link](#))

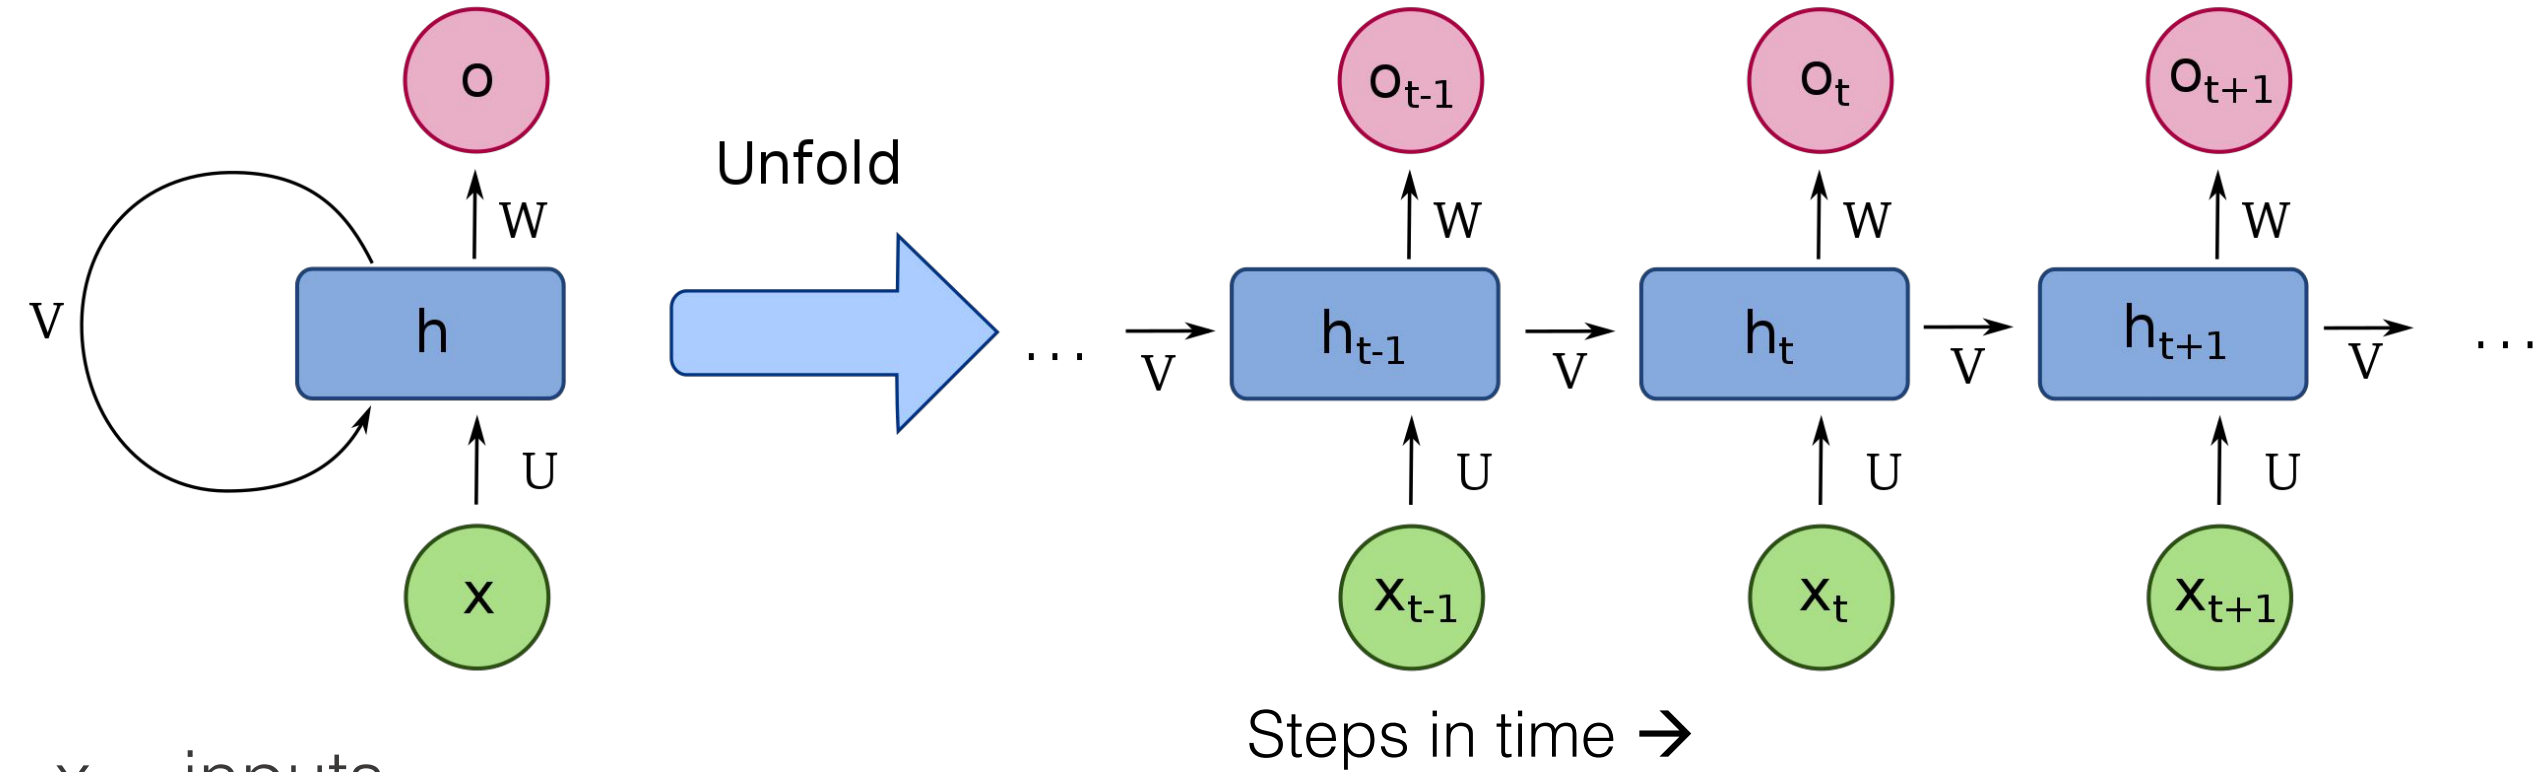
Framework for implementing graphical models, such as neural networks



PyTorch

# KERAS DEMO

# Recurrent Neural Networks



$x$  = inputs

$o$  = outputs

$h$  = hidden layers

$U, V, W$  = model weights

Image from [https://en.wikipedia.org/wiki/Recurrent\\_neural\\_network](https://en.wikipedia.org/wiki/Recurrent_neural_network)



# Generative Adversarial Networks

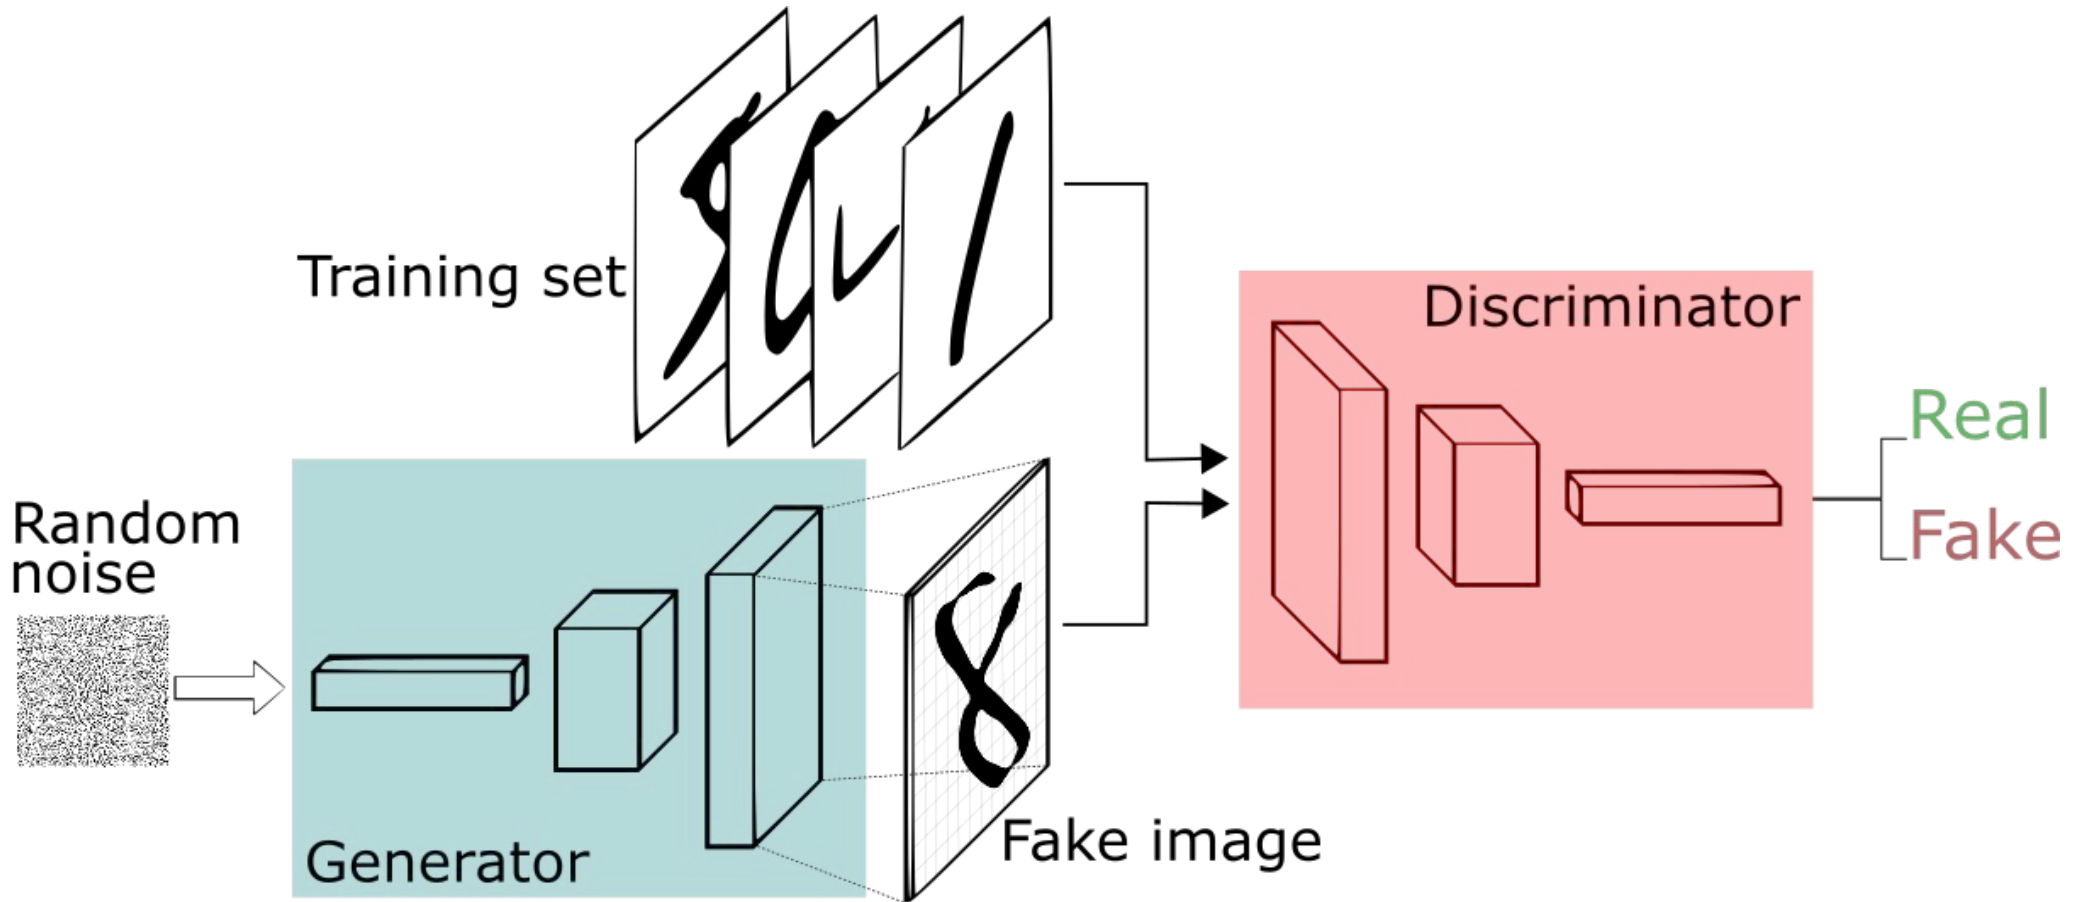


Image from: <https://skymind.ai/wiki/generative-adversarial-network-gan>

# Supervised Learning Techniques

Covered so far

- Linear Regression
- K-Nearest Neighbors
- Perceptron
- Logistic Regression
- Linear Discriminant Analysis
- Quadratic Discriminant Analysis
- Naïve Bayes
- Support Vector Machines
- Decision Trees and Random Forests
- Ensemble methods (bagging, boosting, stacking)
- Neural Networks

Appropriate for:  
● Classification  
● Regression

Can be used with many machine learning techniques